

众核编程基础



MIC

薛巍

2012.06.05


The Challenge of Parallelism

- Programming parallel processors is one of the challenges of our era




NVIDIA Tegra 2 system on a chip (SoC)

- Dual-core ARM Cortex A9.
- Integrated GPU. Lots of DSP.
- 1 GHz.
- 2 single-precision GFLOPs peak (CPUs only)



Nvidia Fermi

- 16 cores, 48-way multithreaded,
- 4-wide Superscalar, dual-issue, 3
- 2-wide SIMD (half-pumped)
- 2 MB (16 x 128 KB) Registers, 1
- MB (16 x 64 KB) L1 cache, 0.75 MB L2 Cache




Tiler Tile64

- 64 processors
- Each tile has L1, L2, can run OS
- 443 billion operations/sec.
- 500-833 MHz
- 50 Gbytes/sec memory bandwidth

Energy Efficient Computing is Key to Performance Growth

At \$1M per MW, energy costs are substantial

- 1 petaflop in 2010 used 3 MW
- 1 exaflop in 2018 would use 100+ MW with “Moore’s Law” scaling




This problem doesn't change if we were to build 1000 1-Petaflop machines instead of 1 Exasflop machine. It affects every university department cluster and cloud data center.

Challenge: New Processor Designs are Needed to Save Energy



Cell phone processor (0.1 Watt, 4 Gflop/s)

Server processor
(100 Watts, 50 Gflop/s)

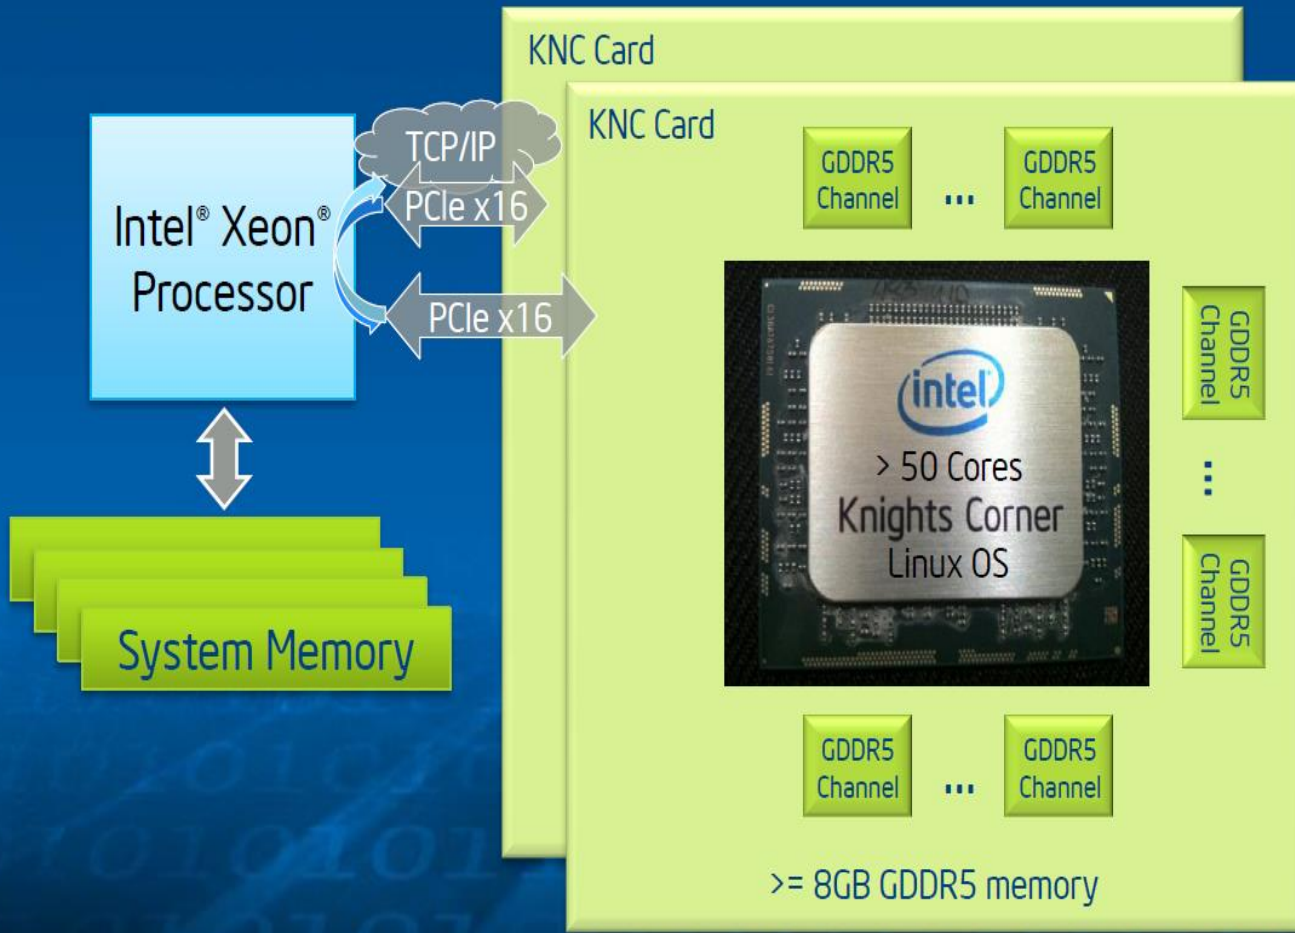


- Server processors have been designed for performance, not energy
 - Graphics processors are 10-100x more efficient
 - Embedded processors are 100-1000x
 - Need manycore chips with thousands of cores

What is MIC

- Intel Many Integrated Core Architecture
 - Targeted at highly parallel HPC workloads
 - Power efficient cores, support for parallelism
 - Cores: less speculation, threads, wider SIMD
 - Scalability: high BW on die interconnect and memory
 - General Purpose Programming Environment
 - Linux
 - Fortran, C/C++
 - x86 memory model, IEEE 754

Knights Corner Coprocessor



Knight Corner Numbering

Intel® Xeon Phi™ coprocessor

7 1 10 P

Brand

Designates the family of products

Performance Shelf

SKU Number

Suffix

Generation

Indication of relative pricing and performance within generation

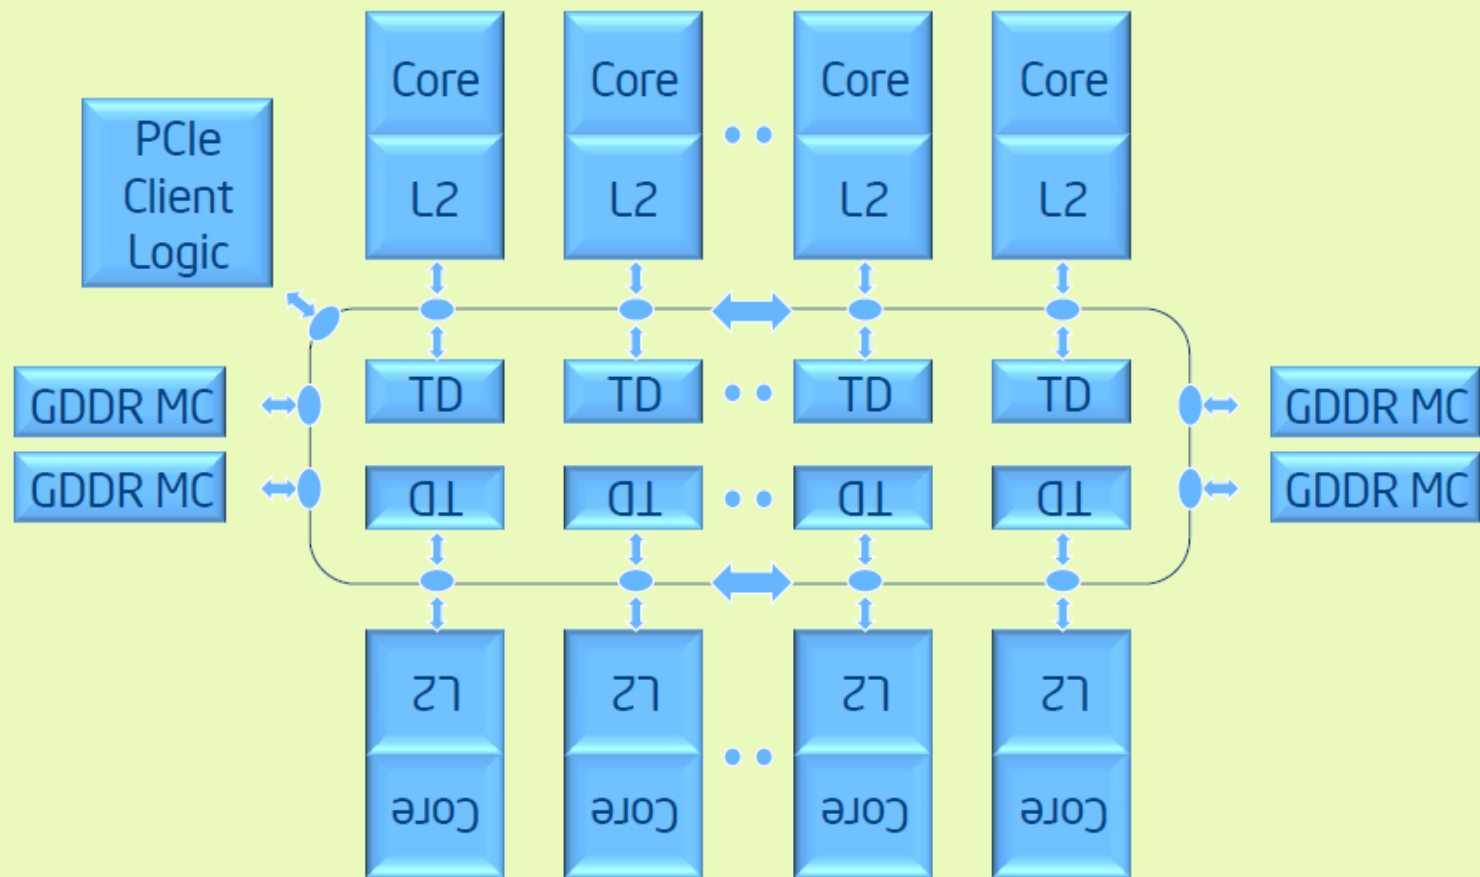
7	Best Performance
5	Best Performance/Watt
3	Best Value

1	Knights Corner
2	Knights Landing

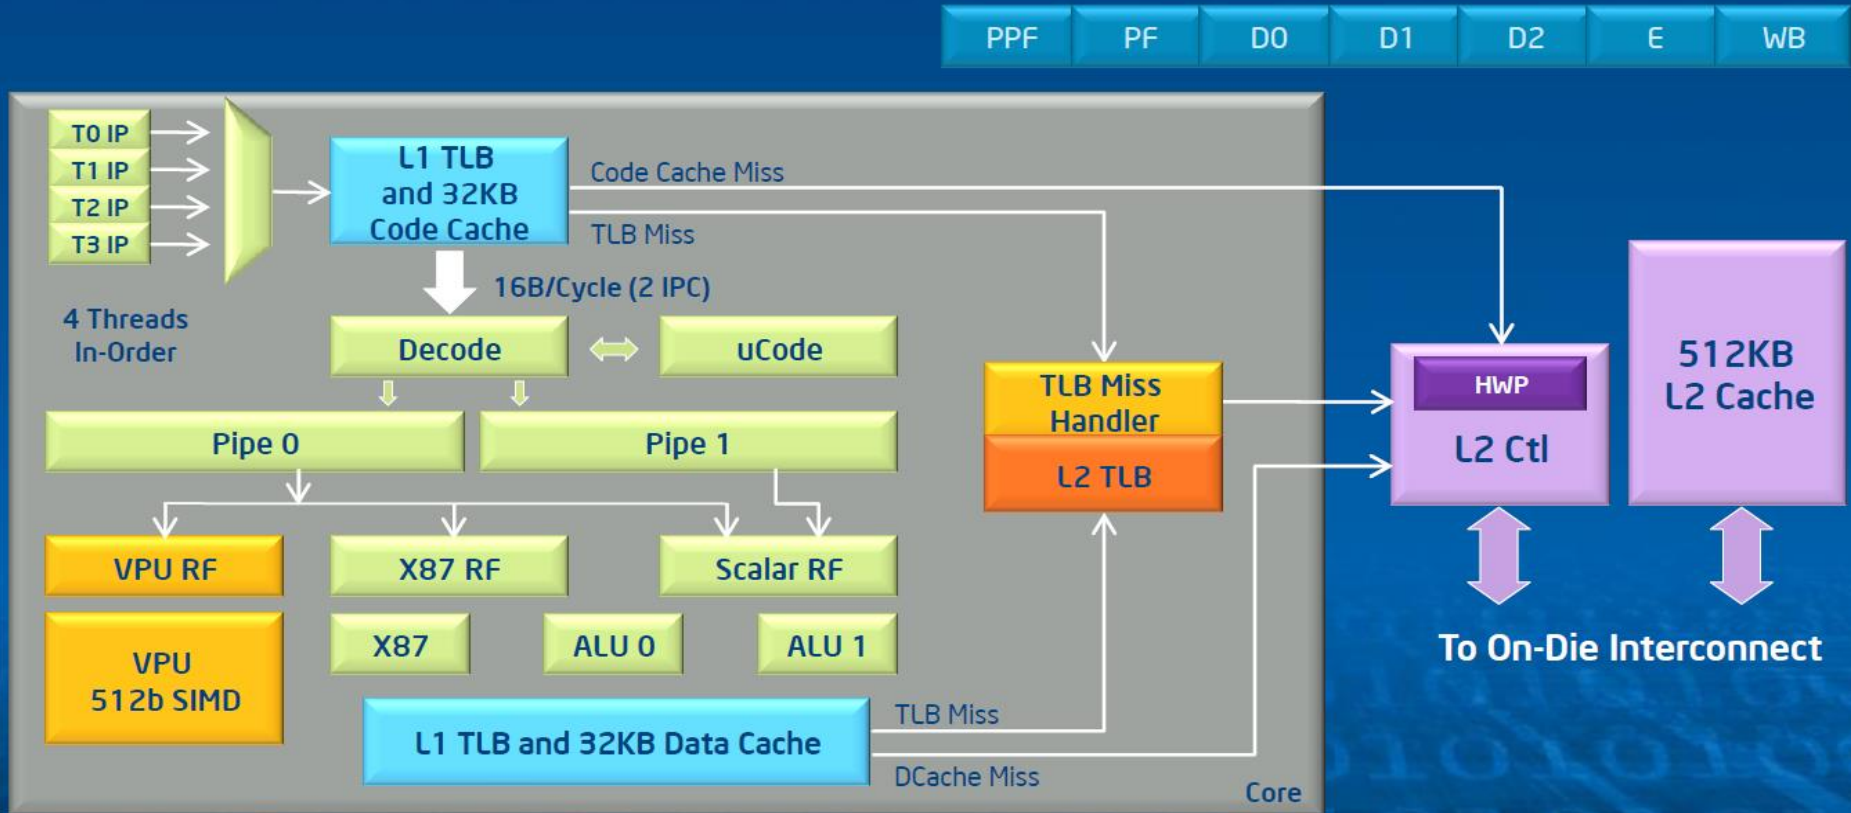
NOTE: Where applicable, may denote form factor, thermal solution, usage model

P	Passive Thermal Solution
A	Active Thermal Solution
X	No Thermal Solution
D	Dense Form Factor

Knights Corner Micro-architecture



Knights Corner Core

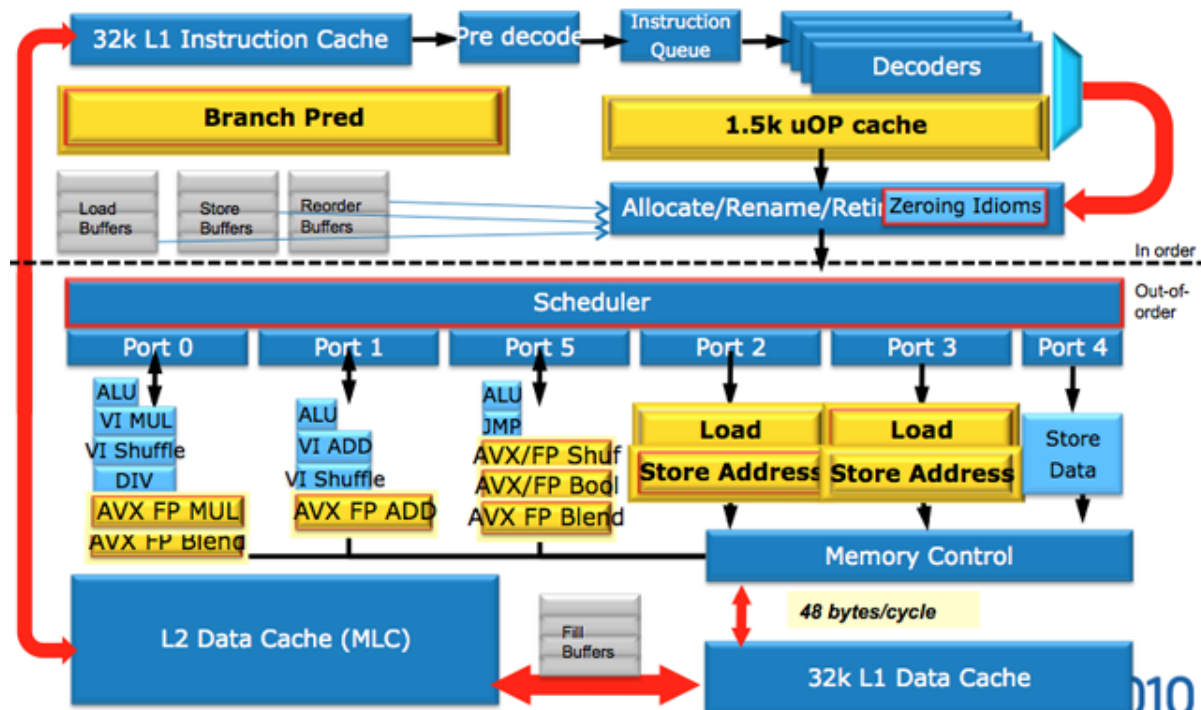


X86 specific logic < 2% of core + L2 area

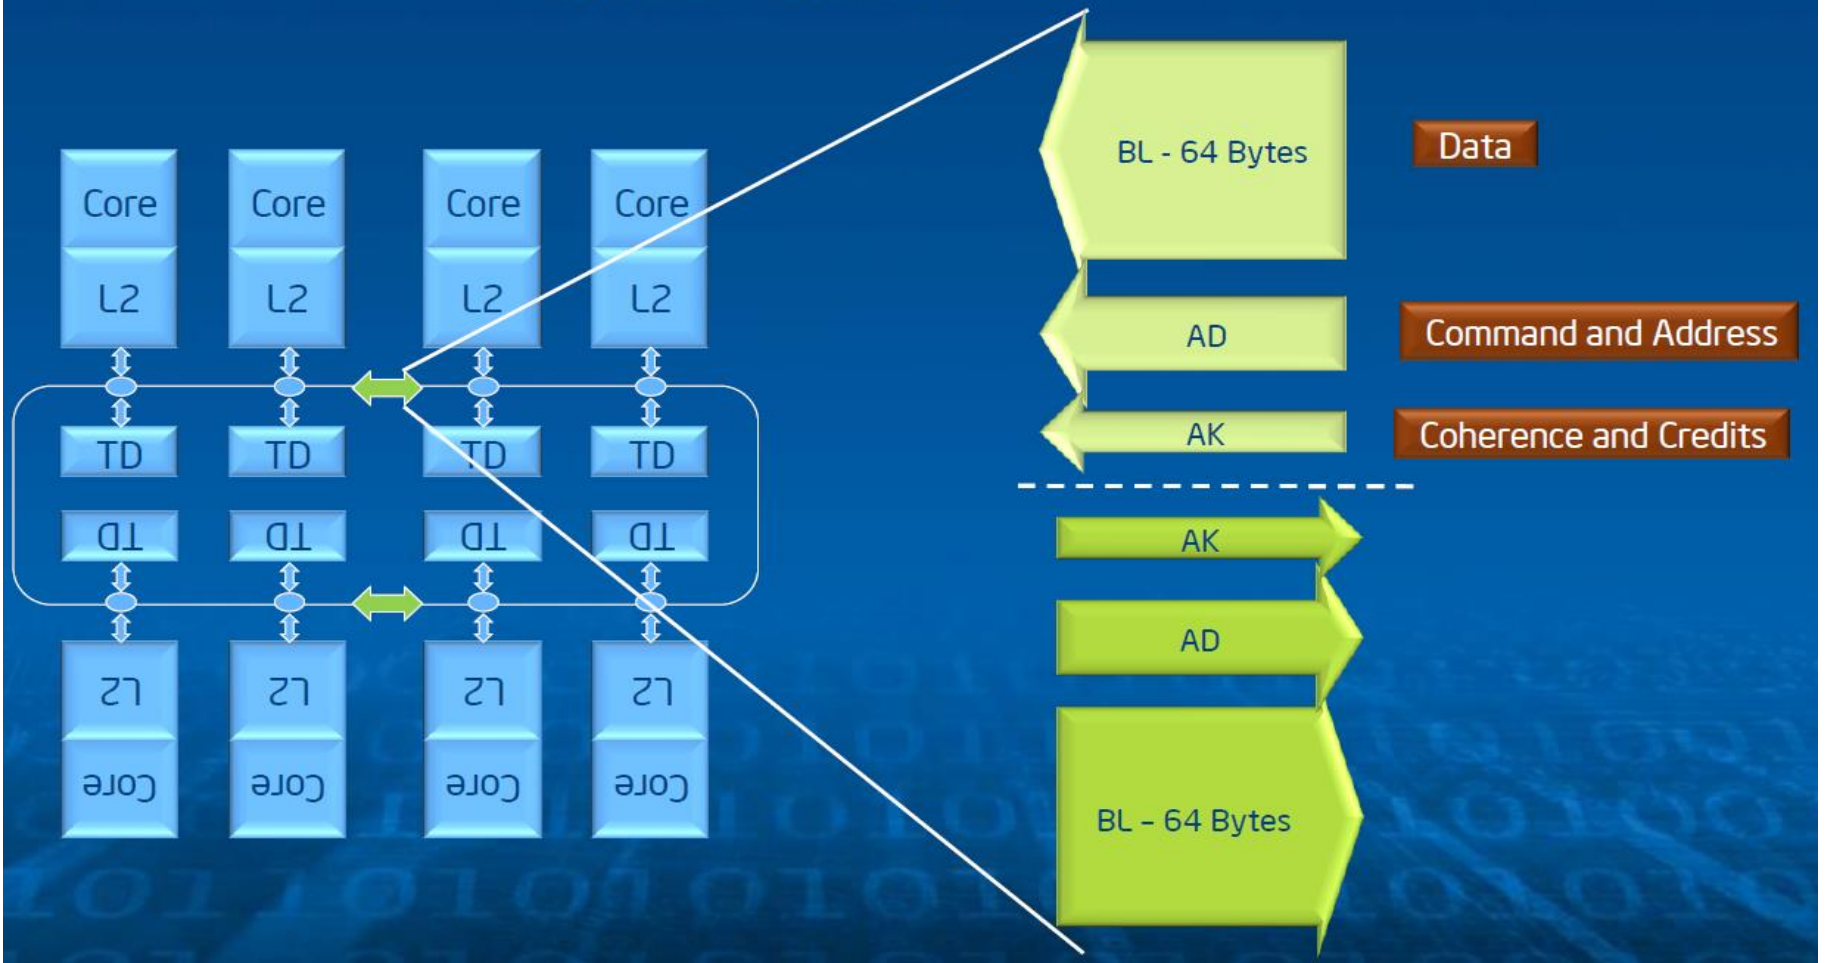
同一线程的指令不能连续发射

Comparing Xeon Phi Core with Xeon Core

Putting it together Sandy Bridge Microarchitecture



Interconnect



Intel Xeon Phi的编程模式

- Native Programming Model
 - entire application running independently on the MIC card
 - evaluate whether your application fits into the card
- Offload Programming Model
 - application decomposed into the host part and the kernel part
 - offload the heavy kernels to the acceleration card
 - similar to the GPU and the FPGA programming model

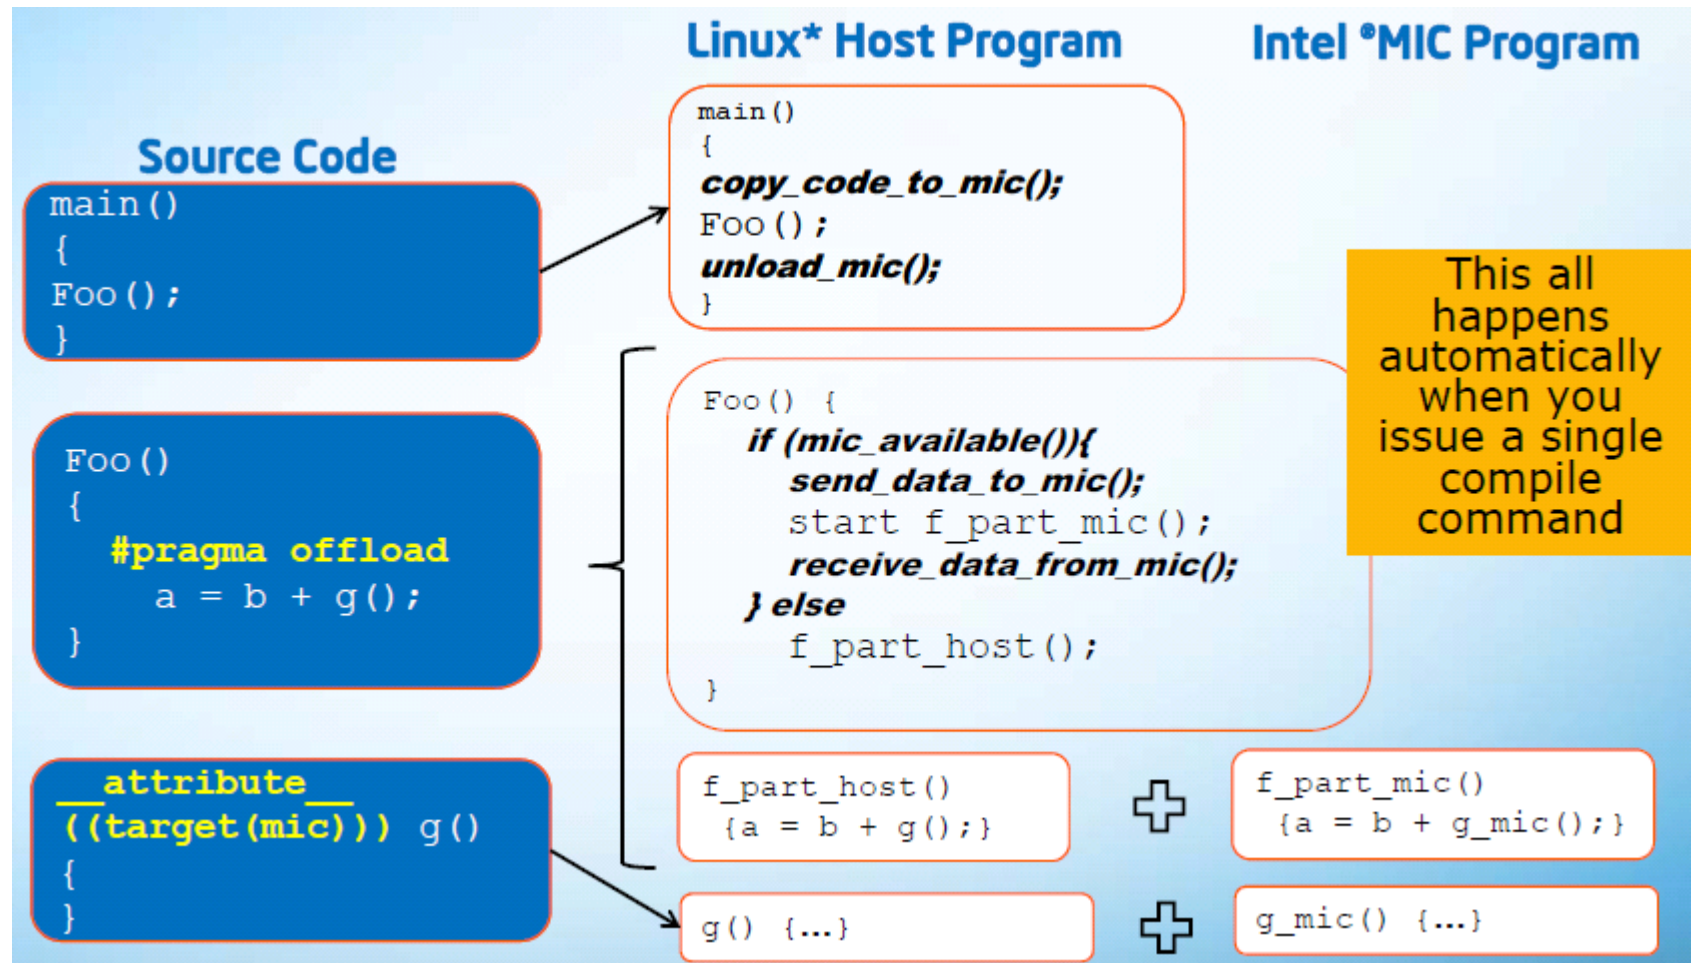
Native Programming Model

- The MIC card is running as an independent node (a simplified Linux)
- write the same code as on CPU
- different compilation and execution methods
 - specific compiling flags: -mmic
 - copy the executable and libraries to the MIC card
- generally provide better performance than the offload mode (20-30%)
- provide the option to use the CPU as the coprocessor

Offload Programming Model

- Add pragmas or keywords to specify the sections that would run on MIC
 - OpenMP:
`#pragma offload target(mic)`
 - Intel Cilk+: `_Cilk_offload_to`
- Automated compilation of different code sections on different architectures
- Running of the code

Compiler Automation for offload code



Offload Data Transfer

- Explicit data transfer:
 - programmer specifies the variables that need to be copied between the host and the card
 - `#pragma offload target(mic) in(data:length(size))`
- Implicit data transfer:
 - programmer marks the variables that need to be shared between the host and the card
 - `_Cilk_shared double foo;`

Example of offload code

- C/C++ Offload Pragma

```
#pragma offload target (mic)
#pragma omp parallel for reduction(+:pi)
for (i=0; i<count; i++) {
    float t = (float)((i+0.5)/count);
    pi += 4.0/(1.0+t*t);
}
pi /= count;
```

- Function Offload Example

```
#pragma offload target(mic)
    in(transa, transb, N, alpha, beta) \
    in(A:length(matrix_elements)) \
    in(B:length(matrix_elements)) \
    inout(C:length(matrix_elements))
    sgemm(&transa, &transb, &N, &N, &N,
&alpha, A, &N, B, &N, &beta, C, &N);
```

- Fortran Offload Directive

```
!dir$ omp offload target(mic)
!$omp parallel do
    do i=1,10
        A(i) = B(i) * C(i)
    enddo
```

- C/C++ Language Extension

```
class _Shared common {
    int data1;
    char *data2;
    class common *next;
    void process();
};


_Shared class common obj1, obj2;
_Cilk_spawn _Offload obj1.process();
_Cilk_spawn          obj2.process();
```

Example of automatic offload

- Possible with advanced Intel Libraries – e.g. MKL

```
void foo() /* Intel® Math Kernel Library */  
{  
    float *A, *B, *C; /* Matrices */  
    sgemm(&transa, &transb, &N, &N, &N, &alpha, A, &N, B, &N, &beta, C, &N);  
}
```

Implicit automatic offloading requires no code changes, simply link with the offload MKL Library



Intel® Xeon® processor



Intel® Xeon Phi™ coprocessor

Heterogeneous Compiler – What is generated

Note that for both techniques, the compiler generates two binaries:

- The host version
 - **includes all functions/variables** in the source code, whether marked `#pragma offload`, `__attribute__((target(mic)))`, `_Cilk_shared`, `_Cilk_offload`, or not
- The coprocessor version
 - **includes only functions/variables** marked `#pragma offload`, `__attribute__((target(mic)))`, `_Cilk_offload`, or `_Cilk_shared` in the source code

Heterogeneous Compiler – Command-line option

- “-openmp” is automatically set in offload
- Most command line arguments set for the host are set for the coprocessor build
 - Unless overridden by `-offload-option,mic,...`
- Setting up the compiler build environment

```
CSH: source /opt/intel/composerxe_mic/bin/compilervars.csh
intel64
SH:   source /opt/intel/composerxe_mic/bin/compilervars.sh
intel64
```


Heterogeneous Compiler – Command-line option

Offload-specific arguments to the Intel® Compiler:

- Ignore language constructs for offloading (by default the compiler recognizes language constructs for offloading if they are specified):
`-no-offload`
- Produce a report of offload data transfers at compile time (not runtime)
`-opt-report-phase:offload`
- Specify options to be used for Intel® MIC Architecture
`-offload-option,mic,tool,"option-list"`

Example:

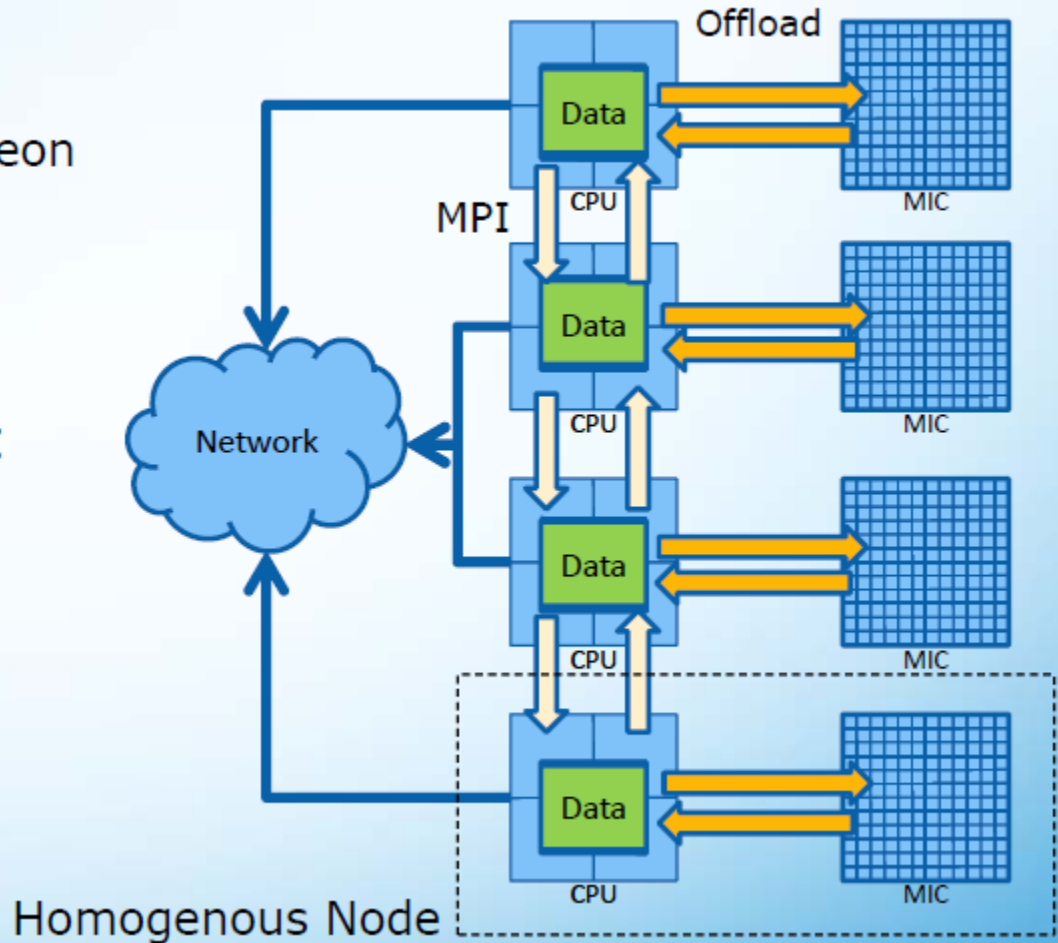
```
Icc -I/my_dir/include -DMY_DEFINE=10 -offload-option,mic,compiler, "-  
I/my_dir/mic/include -DMY_DEFINE=20" -offload-option,mic,ld, "-  
L/my_dir/mic/lib" hello.c
```

Passes "-I/my_dir/mic/include -I/my_dir/include -DMY_DEFINE=10 -
DMY_DEFINE=20" to the offload compiler

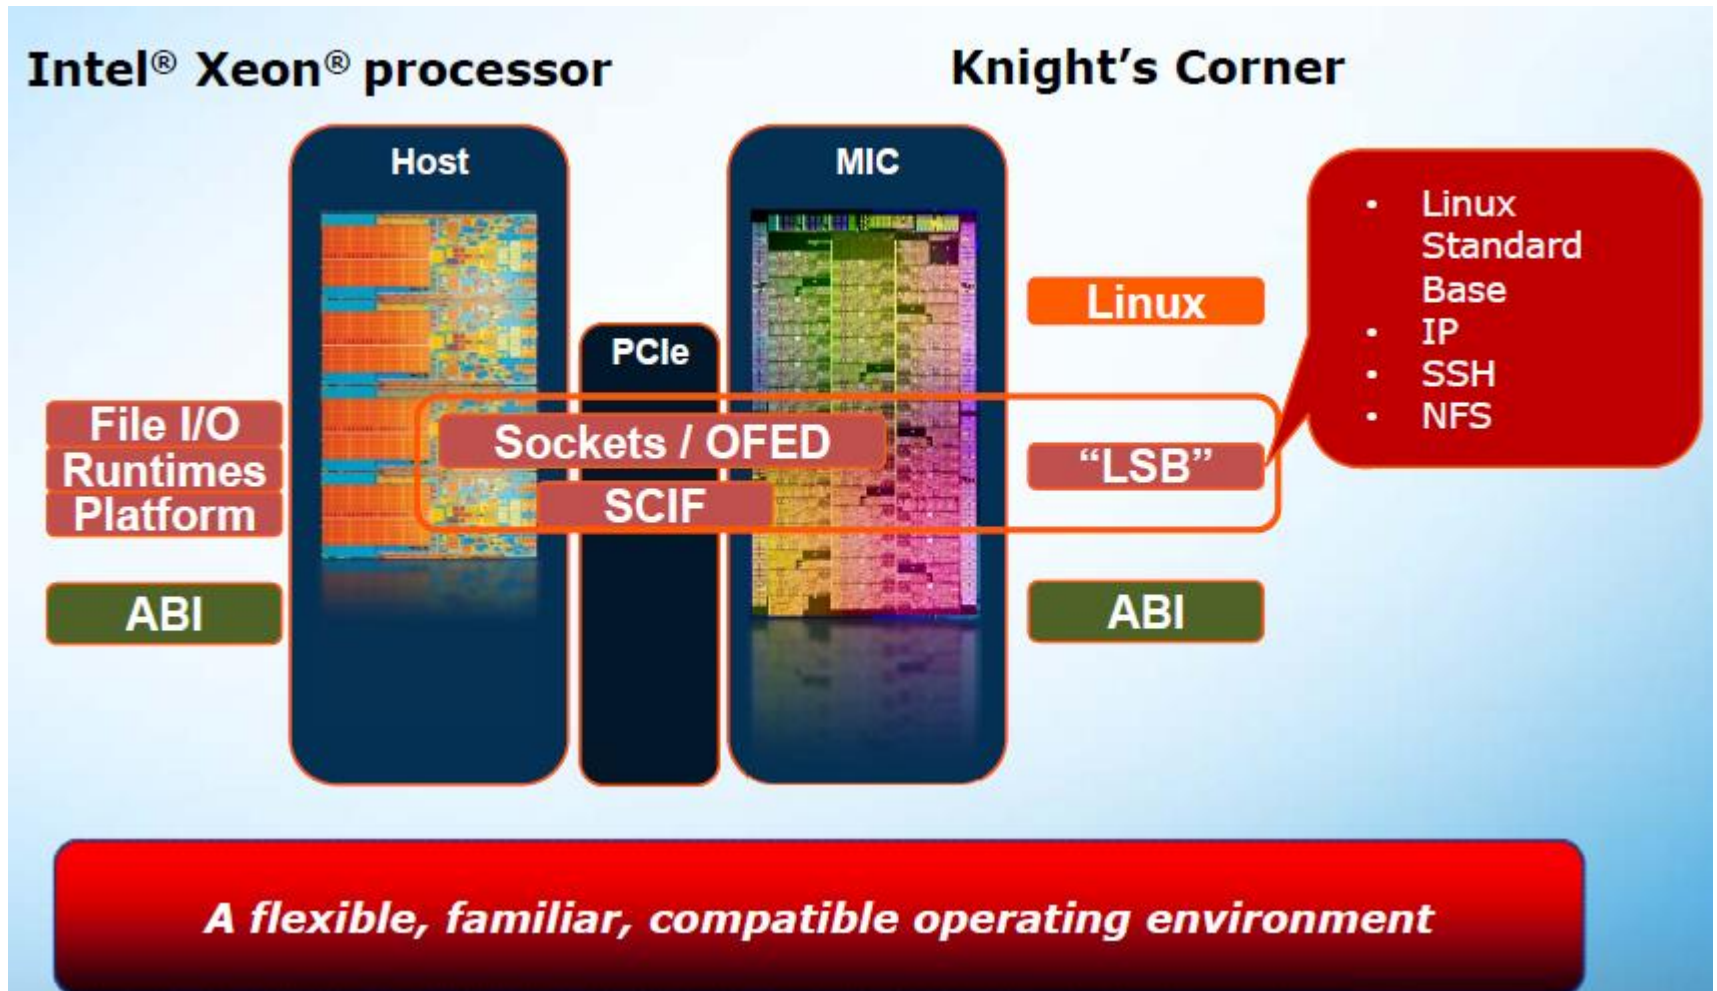
Passes "-L/my_dir/mic/lib -L/my_dir/lib" to the offload ld

Offload + MPI

- MPI ranks on Intel® Xeon® processors (only)
- All messages into/out of Xeon processors
- Offload models used to accelerate MPI ranks
- TBB, OpenMP, Cilk Plus, Pthreads within Intel® MIC
- Homogenous network of hybrid nodes



Operating Environment




Key points for MIC

- Choose the right Xeon-centric or MIC-centric model for your application
- Vectorize your application
 - Use Intel's vectorizing compiler
- Parallelize your application
 - With MPI (or other multiprocess model)
 - With threads (via Pthreads, TBB, Cilk Plus, OpenMP, etc.)
- Consider standard optimization techniques such as tiling, overlapping computation and communication, etc.

Computation of PI

Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

Serial Version

```
static long num_steps = 100000;
double step;
void main ()
{
    int i;  double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;

    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```


Experiments

- Pi
 - Serial version
 - OpenMP CPU version
 - OpenMP MIC version (Native and Offloading)
- 期权定价问题的step5

Guide for experiment

- Please login 166.111.68.173 with putty
 - user:test
 - passwd:test@tsinghua
- Go to your dir in home dir
 - cd xuew
- Copy cases to your dir
 - cp -pr /tmp/lec11/MonteCarlo/ .
 - cp -pr ~/pi.cpp .

课程回顾

- 本课程的目的
 - 偏实践类课程
 - 实实在在体验并行软件开发过程
 - 并行编程 — 实现
 - 并行程序性能分析与优化 — 分析优化
 - 并行算法/程序设计 — 设计
 - 并行计算工具 — 组合
 - 一些对于计算机的感性认识
- 预期效果
 - 在今后采用并行计算方法提升程序性能
 - 能编写不复杂的并行程序
 - 知道使用并行工具提高并行程序实现效率
 - 初步了解并行计算程序设计和优化的过程
- 进一步深入学习可以首先参考：
<http://www.cs.berkeley.edu/~demmel/cs267>

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1K bytes with Zip	10,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

计算机系并行计算相关课程

计算机系统性能测试与
分析

高级编译与优化技术

大规模科学计算 计算科学与工程中的
并行编程技术

网络存储技术

网络计算

并行计算

高等计算机体系结构

高性能计算前沿
技术探讨

高性能计算
导论 并行计算基础

高性能计算平台
与工具

Using Language Extensions for Intel® MIC

Simple Offload Extensions with the Intel® Compilers

	C/C++ Syntax	Semantics
New offload pragma	<code>#pragma offload (clauses)</code>	Execute next statement on target (which could be an OpenMP* parallel construct)
Place function on target	<code>__declspec (target (x))</code>	Compile function for host and target
Place on data target	<code>__declspec (target(MIC)) float array [8000];</code>	Two arrays are created, one on the host and one on the Intel® Xeon Phi™ coprocessor

	Fortran Syntax	Semantics
New offload directive	<code>!dir\$ omp offload <clauses></code>	Execute next OpenMP* parallel construct on target
Place function on target	<code>!dir\$ attributes offload:<x> :: <rtn-name></code>	Compile function for host and target

Using Language Extensions (contd.)

Variables restricted to scalars, arrays and pointers to scalars/arrays, structs (which can be bit-wise copied)

What	Syntax	Semantics
Target specification	target (name)	Where to run construct
Inputs	in (var-list modifiers _{opt})	Copy CPU to target
Outputs	out (var-list modifiers _{opt})	Copy target to CPU
Inputs & outputs	inout (var-list modifiers _{opt})	Copy both ways
Non-copied data	nocopy (var-list modifiers _{opt})	Data is local to target
Modifiers		
Specify pointer length	length (element-count-expr)	Copy that many pointer elements
Control pointer memory allocation	alloc_if (condition)	Allocate new block of memory for pointer if condition is TRUE
Control freeing of pointer memory	free_if (condition)	Free memory used for pointer if condition is TRUE